

NumPy

1. NumPy Basic Ops

```
In [ ]: import numpy as np  
import pandas as pd
```

-> Create a 1D Array from 1 to 20.

```
In [186... array_1d = np.arange(1, 21)  
print("1D Array from 1 to 20:\n", array_1d)
```

1D Array from 1 to 20:
[1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20]

-> Create a 3×4 matrix of ones and reshape it to 4×3

```
In [189... mat_1 = np.ones((3,4))  
print("\n 3x4 Matrix of Ones: \n", mat_1)  
  
# Reshaping It to 4x3  
  
reshape = mat_1.reshape((4,3))  
print("\n 4x3 Reshaped Materix of Ones: \n", reshape)
```

3x4 Matrix of Ones:
[[1. 1. 1. 1.]
[1. 1. 1. 1.]
[1. 1. 1. 1.]]

4x3 Reshaped Materix of Ones:
[[1. 1. 1.]
[1. 1. 1.]
[1. 1. 1.]
[1. 1. 1.]]

-> Create a 5×5 identity matrix

```
In [194... i5 = np.eye(5)  
print("\n 5x5 Identity Matrix: \n", i5)
```

5x5 Identity Matrix:
[[1. 0. 0. 0. 0.]
[0. 1. 0. 0. 0.]
[0. 0. 1. 0. 0.]
[0. 0. 0. 1. 0.]
[0. 0. 0. 0. 1.]]

-> Generate 15 equally spaced numbers between 5 and 50.

In [199...]

```
e_nums = np.linspace(5,50,15)
print("\n 15 equally spaced nums between 5 and 15: \n", e_nums)
```

```
15 equally spaced nums between 5 and 15:
[ 5.          8.21428571 11.42857143 14.64285714 17.85714286 21.07142857
24.28571429 27.5          30.71428571 33.92857143 37.14285714 40.35714286
43.57142857 46.78571429 50.          ]
```

-> Generate a 4×4 matrix of random integers between 1 and 100.

In [201...]

```
r_nums = np.random.randint(1,101,size=(4,4))
print("\n 4x4 matrix of Random int between 1 & 100 \n",r_nums)
```

```
4x4 matrix of Random int between 1 & 100
[[48 17 75 76]
[47 94 95 48]
[36 68 89 44]
[93 58 61 56]]
```

2. NumPy Array Ops

-> Create a 3×3 matrix of random integers between 1 and 100

In [205...]

```
r_mat = np.random.randint(1,100,size=(3,3))
print("\n 3x3 Matrix of random nums \n",r_mat)
```

```
3x3 Matrix of random nums
[[34 34 99]
[14 14 29]
[37 85  1]]
```

-> Extract: First row, Second column, Center element.

In [208...]

```
fr = r_mat[0]
print("\n First Row: \n",fr)

sc = r_mat[:,1]
print("\n Second Column \n",sc)

se = r_mat[1,1]
print("\n Center Element \n",se)
```

First Row:

[34 34 99]

Second Column

[34 14 85]

Center Element

14

-> Replace all values greater than 50 in a matrix with 999.

```
In [211...]
r_mat_c = r_mat.copy()
r_mat_c[r_mat_c > 50] = 999
print("\n Replaced Matrix \n",r_mat_c)
```

Replaced Matrix

[[34 34 999]
[14 14 29]
[37 999 1]]

-> Multiply a 1D array of size 5 (random integers between 1 to 10) by 10 using broadcasting.

```
In [214...]
arr_1d = np.random.randint(1,11,size=5)
multi_arr = arr_1d * 10
print("\n Original 1D Array: \n",arr_1d)
print("\n Multiplied Array: \n",multi_arr)
```

Original 1D Array:

[10 9 3 8 6]

Multiplied Array:

[100 90 30 80 60]

3. Math & Stats Ops

-> Create a 3×3 matrix of random integers between 1 and 100

```
In [225...]
r_mat = np.random.randint(1,100,size=(3,3))
print("\n 3x3 Matrix of random nums \n",r_mat)
```

3x3 Matrix of random nums

[[24 1 95]
[79 17 61]
[99 93 46]]

-> Compute sum, mean, median, std, var, min, and max of the above array.

```
In [230...]
print("\nStatistics:")
print("Sum:", np.sum(r_mat))
print("Mean:", np.mean(r_mat))
```

```

print("Median:", np.median(r_mat))
print("Standard Deviation:", np.std(r_mat))
print("Variance:", np.var(r_mat))
print("Minimum:", np.min(r_mat))
print("Maximum:", np.max(r_mat))

```

Statistics:
 Sum: 515
 Mean: 57.22222222222222
 Median: 61.0
 Standard Deviation: 34.880041341010724
 Variance: 1216.6172839506173
 Minimum: 1
 Maximum: 99

-> Normalize a 1D array of size 5 (random integers between 1 to 10) to scale values between 0 and 1.

In [234...]

```

arr_1d = np.random.randint(1,11,size=5)
print("\n 1D Array: \n",arr_1d)
min = np.min(arr_1d)
max = np.max(arr_1d)
norm_arr = ((arr_1d - min) / (max - min) )
print("\n Normalized Array: \n",norm_arr)

```

1D Array:
 [8 10 5 1 4]
 Normalized Array:
 [0.77777778 1. 0.44444444 0. 0.33333333]

4. NumPy Matrix Ops & Algebra

-> Generate Following two NumPy matrix import numpy as np

A = ([4, 2], [1, 3])

B = ([2, 0], [1, 5])

In [240...]

```

A = np.array( [[4, 2], [1, 3]] )
B = np.array( [[2, 0], [1, 5]] )

print("\n A: \n",A)
print("\n B: \n",B)

```

```
A:  
[[4 2]  
[1 3]]
```

```
B:  
[[2 0]  
[1 5]]
```

-> Matrix Multiplication

```
In [252... mat_multi = np.matmul(A,B)  
print("\n Multiplication of A & B: \n",mat_multi)
```

```
Multiplication of A & B:  
[[10 10]  
[ 5 15]]
```

-> dot product of A and B.

```
In [254... dot = np.dot(A,B)  
print("\n dot product of A & B.: \n",dot)
```

```
dot product of A & B.:  
[[10 10]  
[ 5 15]]
```

-> Element wise addition/ subtraction/ multiplications/ division of A and B

```
In [250... add = A + B  
sub = A - B  
multi = A * B  
div = A / B  
  
print("\nElement-wise Addition:\n", add)  
print("Element-wise Subtraction:\n", sub)  
print("Element-wise Multiplication:\n", multi)  
print("Element-wise Division:\n", div)
```

```
Element-wise Addition:  
[[6 2]  
[2 8]]  
Element-wise Subtraction:  
[[ 2  2]  
[ 0 -2]]  
Element-wise Multiplication:  
[[ 8  0]  
[ 1 15]]  
Element-wise Division:  
[[2. inf]  
[1.  0.6]]
```

```
C:\Users\rajva\AppData\Local\Temp\ipykernel_18728\3905969198.py:4: RuntimeWarning: divide by zero encountered in divide  
div = A / B
```

-> Transpose of A

```
In [277...]: trans_A = A.T
print("\n Transpose of A: \n", trans_A)
```

Transpose of A:
[[4 1]
[2 3]]

-> Determinant of A

```
In [271...]: det_A = np.linalg.det(A)
print("\n Determinant of A: \n", det_A)
```

Determinant of A:
10.00000000000002

-> Inverse of A

```
In [269...]: if det_A != 0:
    inv_A = np.linalg.inv(A)
    print("\nInverse of A:\n", inv_A)
else:
    print("\nMatrix A cannot be inverted.")
```

Inverse of A:
[[0.3 -0.2]
[-0.1 0.4]]

-> Eigenvalues and Eigenvectors of A

```
In [264...]: eigenvalues= np.linalg.eig(A)
eigenvectors = np.linalg.eig(A)
print("\nEigenvalues of A:\n", eigenvalues)
print("\nEigenvectors of A:\n", eigenvectors)
```

Eigenvalues of A:
EigResult(eigenvalues=array([5., 2.]), eigenvectors=array([[0.89442719, -0.70710678],
[0.4472136 , 0.70710678]]))

Eigenvectors of A:
EigResult(eigenvalues=array([5., 2.]), eigenvectors=array([[0.89442719, -0.70710678],
[0.4472136 , 0.70710678]]))

-> Solve System of Equations:

$$2x + y = 8$$

$$3x + 4y = 18$$

```
In [261...]: coeff = np.array([[2,1],[3,4]])
cons = np.array([8,18])

sol = np.linalg.solve(coeff, cons)
print("\n Solution of the Equations x & y : ",sol)
```

Solution of the Equations x & y : [2.8 2.4]

Pandas Exercise

1. Series & DataFrame Basics

-> Given the following list of marks:

[78, 85, 92, 70, 66]

Create a Pandas Series and assign the following student names as indices:

['Amit', 'Bhavna', 'Chetan', 'Divya', 'Esha']

Display the Series

```
In [ ]: marks = [78, 85, 92, 70, 66]
names = ['Amit', 'Bhavna', 'Chetan', 'Divya', 'Esha']

marks_series = pd.Series(data=marks, index=names)
print("\nPandas Series : \n",marks_series)
```

-> Create a Pandas DataFrame and display:

The full DataFrame

The column names

The shape of the DataFrame

```
In [285...]: # Create DataFrame
data = {
    'Name': ['Amit', 'Bhavna', 'Chetan', 'Divya', 'Esha'],
    'Gender': ['Male', 'Female', 'Male', 'Female', 'Female'],
    'Math': [78, 85, 92, 70, 66],
    'Science': [88, 79, 95, 72, 60]
}

df = pd.DataFrame(data)

print("\nFull DataFrame:\n")
print(df)
```

```
# Display column names
print("\nColumn Names:\n")
print(df.columns.tolist())

# Display shape (rows, columns)
print("\nShape of the DataFrame:\n")
print(df.shape)
```

Full DataFrame:

	Name	Gender	Math	Science
0	Amit	Male	78	88
1	Bhavna	Female	85	79
2	Chetan	Male	92	95
3	Divya	Female	70	72
4	Esha	Female	66	60

Column Names:

```
[ 'Name', 'Gender', 'Math', 'Science' ]
```

Shape of the DataFrame:

(5, 4)

2. Data Exploration

-> Load the dataset from this url
["https://archive.ics.uci.edu/ml/machine-learningdatabases/autos/imports-85.data"](https://archive.ics.uci.edu/ml/machine-learningdatabases/autos/imports-85.data)

```
In [7]: import pandas as pd  
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.d  
df = pd.read_csv(url, header=None)  
print(df.head())
```

	0	1	2	3	4	5	6	7	8	9	...	\	
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...		
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...		
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...		
3	2	164		audi	gas	std	four		sedan	fwd	front	99.8	...
4	2	164		audi	gas	std	four		sedan	4wd	front	99.4	...
	16	17	18	19	20	21	22	23	24	25			
0	130	mpfi	3.47	2.68	9.0	111	5000	21	27	13495			
1	130	mpfi	3.47	2.68	9.0	111	5000	21	27	16500			
2	152	mpfi	2.68	3.47	9.0	154	5000	19	26	16500			
3	109	mpfi	3.19	3.40	10.0	102	5500	24	30	13950			
4	136	mpfi	3.19	3.40	8.0	115	5500	18	22	17450			

[5 rows x 26 columns]

-> Assign names of columns:

```
["symboling","normalized-losses","make","fuel-type","aspiration", "num-ofdoors","body-style", "drive-wheels","engine-location","wheelbase","length","width","height","curb-weight","engine-type","num-of-cylinders","engine-size","fuel-system","bore","stroke","compression-ratio","horsepower","peak-rpm","city-mpg","highway-mpg","price"]
```

In [296...]

```
df.columns = ["symboling","normalized-losses","make","fuel-type","aspiration", "num-of-doors", "body-style", "drive-wheels", "engine-location", "wheel-base", "...", "engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower", "peak-rpm", "city-mpg", "highway-mpg", "price"]
print(df.head())
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg	price
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21	27	13495
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21	27	16500
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154	5000	19	26	16500
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102	5500	24	30	13950
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115	5500	18	22	17450

[5 rows x 26 columns]

-> Display .shape, .columns, .info(), and .describe()

In [300...]

```
print(df.shape)
print(df.columns)
print(df.info())
print(df.describe())
```

(205, 26)

```

Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
       'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
       'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
       'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
       'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
       'highway-mpg', 'price'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column            Non-Null Count Dtype  
 --- 
 0   symboling         205 non-null   int64  
 1   normalized-losses 205 non-null   object  
 2   make              205 non-null   object  
 3   fuel-type          205 non-null   object  
 4   aspiration         205 non-null   object  
 5   num-of-doors       205 non-null   object  
 6   body-style          205 non-null   object  
 7   drive-wheels        205 non-null   object  
 8   engine-location     205 non-null   object  
 9   wheel-base          205 non-null   float64 
 10  length             205 non-null   float64 
 11  width              205 non-null   float64 
 12  height              205 non-null   float64 
 13  curb-weight         205 non-null   int64  
 14  engine-type         205 non-null   object  
 15  num-of-cylinders    205 non-null   object  
 16  engine-size          205 non-null   int64  
 17  fuel-system          205 non-null   object  
 18  bore                205 non-null   object  
 19  stroke              205 non-null   object  
 20  compression-ratio    205 non-null   float64 
 21  horsepower           205 non-null   object  
 22  peak-rpm             205 non-null   object  
 23  city-mpg             205 non-null   int64  
 24  highway-mpg          205 non-null   int64  
 25  price               205 non-null   object  
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB
None
      symboling  wheel-base    length      width      height \
count  205.000000  205.000000  205.000000  205.000000  205.000000
mean    0.834146   98.756585  174.049268   65.907805   53.724878
std     1.245307   6.021776  12.337289   2.145204   2.443522
min    -2.000000   86.600000  141.100000   60.300000   47.800000
25%    0.000000   94.500000  166.300000   64.100000   52.000000
50%    1.000000   97.000000  173.200000   65.500000   54.100000
75%    2.000000  102.400000  183.100000   66.900000   55.500000
max    3.000000  120.900000  208.100000   72.300000   59.800000

      curb-weight  engine-size  compression-ratio  city-mpg  highway-mpg
count  205.000000  205.000000      205.000000  205.000000  205.000000
mean   2555.565854  126.907317      10.142537  25.219512  30.751220
std    520.680204   41.642693      3.972040   6.542142  6.886443

```

min	1488.000000	61.000000	7.000000	13.000000	16.000000
25%	2145.000000	97.000000	8.600000	19.000000	25.000000
50%	2414.000000	120.000000	9.000000	24.000000	30.000000
75%	2935.000000	141.000000	9.400000	30.000000	34.000000
max	4066.000000	326.000000	23.000000	49.000000	54.000000

-> Display only "width", "height", "curb-weight", "engine-type" columns.

In [303...]

```
print(df[["width", "height", "curb-weight", "engine-type"]])
```

```
   width  height  curb-weight engine-type
0     64.1    48.8       2548      dohc
1     64.1    48.8       2548      dohc
2     65.5    52.4       2823      ohcv
3     66.2    54.3       2337      ohc
4     66.4    54.3       2824      ohc
..     ...
200    68.9    55.5       2952      ohc
201    68.8    55.5       3049      ohc
202    68.9    55.5       3012      ohcv
203    68.9    55.5       3217      ohc
204    68.9    55.5       3062      ohc
```

[205 rows x 4 columns]

-> Display car details which have num-of-doors = four

In [307...]

```
print(df[df["num-of-doors"] == "four"])
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	
6	1	158	audi	gas	std	four	
7	1	?	audi	gas	std	four	
8	1	158	audi	gas	turbo	four	
..	
200	-1	95	volvo	gas	std	four	
201	-1	95	volvo	gas	turbo	four	
202	-1	95	volvo	gas	std	four	
203	-1	95	volvo	diesel	turbo	four	
204	-1	95	volvo	gas	turbo	four	
	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	
6	sedan	fwd	front	105.8	...	136	
7	wagon	fwd	front	105.8	...	136	
8	sedan	fwd	front	105.8	...	131	
..	
200	sedan	rwd	front	109.1	...	141	
201	sedan	rwd	front	109.1	...	141	
202	sedan	rwd	front	109.1	...	173	
203	sedan	rwd	front	109.1	...	145	
204	sedan	rwd	front	109.1	...	141	
	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	\
3	mpfi	3.19	3.40	10.0	102	5500	
4	mpfi	3.19	3.40	8.0	115	5500	
6	mpfi	3.19	3.40	8.5	110	5500	
7	mpfi	3.19	3.40	8.5	110	5500	
8	mpfi	3.13	3.40	8.3	140	5500	
..	
200	mpfi	3.78	3.15	9.5	114	5400	
201	mpfi	3.78	3.15	8.7	160	5300	
202	mpfi	3.58	2.87	8.8	134	5500	
203	idi	3.01	3.40	23.0	106	4800	
204	mpfi	3.78	3.15	9.5	114	5400	
	city-mpg	highway-mpg	price				
3	24	30	13950				
4	18	22	17450				
6	19	25	17710				
7	19	25	18920				
8	17	20	23875				
..				
200	23	28	16845				
201	19	25	19045				
202	18	23	21485				
203	26	27	22470				
204	19	25	22625				

[114 rows x 26 columns]

3. Missing Value Handlling

-> Missing value is represented by '?' in this dataset. Replace it with NULL.

```
In [311... df.replace('?', None , inplace=True)
print(df.head())
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	3	None	alfa-romero	gas	std	two	
1	3	None	alfa-romero	gas	std	two	
2	1	None	alfa-romero	gas	std	two	
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

-> Check how many missing values are there in each attribute.

```
In [316... print(df.isnull().sum())
```

```
symboling          0
normalized-losses 41
make              0
fuel-type         0
aspiration        0
num-of-doors      2
body-style        0
drive-wheels      0
engine-location   0
wheel-base        0
length             0
width              0
height             0
curb-weight        0
engine-type        0
num-of-cylinders  0
engine-size        0
fuel-system        0
bore               4
stroke             4
compression-ratio 0
horsepower         2
peak-rpm            2
city-mpg            0
highway-mpg         0
price              4
dtype: int64
```

-> Replace missing values of "normalized-losses", "stroke", "bore", "horsepower" with mean.

```
In [26]: for col in ["normalized-losses", "stroke", "bore", "horsepower"]:
    df[col].fillna(df[col].mean(), inplace=True)
print(df.head())
```

```

symboling normalized-losses      make fuel-type aspiration \
0            3          122.0 alfa-romero     gas      std
1            3          122.0 alfa-romero     gas      std
2            1          122.0 alfa-romero     gas      std
3            2          164.0    audi     gas      std
4            2          164.0    audi     gas      std

  num-of-doors body-style drive-wheels engine-location wheelbase ... \
0        two   convertible           rwd       front     88.6 ...
1        two   convertible           rwd       front     88.6 ...
2        two   hatchback            rwd       front     94.5 ...
3       four    sedan              fwd       front     99.8 ...
4       four    sedan              4wd       front     99.4 ...

  engine-size fuel-system bore stroke compression-ratio horsepower \
0         130      mpfi  3.47   2.68          9.0     111.0
1         130      mpfi  3.47   2.68          9.0     111.0
2         152      mpfi  2.68   3.47          9.0     154.0
3         109      mpfi  3.19   3.40         10.0     102.0
4         136      mpfi  3.19   3.40          8.0     115.0

  peak-rpm city-mpg highway-mpg   price
0  5000.0       21        27 13495.0
1  5000.0       21        27 16500.0
2  5000.0       19        26 16500.0
3  5500.0       24        30 13950.0
4  5500.0       18        22 17450.0

```

[5 rows x 26 columns]

C:\Users\rajva\AppData\Local\Temp\ipykernel_19968\497043668.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an `inplace` method.

The behavior will change in pandas 3.0. This `inplace` method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation `inplace` on the original object.

```
df[col].fillna(df[col].mean(), inplace=True)
```

-> Drop all the rows which has missing value in attribute "price".

In [31]: `df.dropna(subset=["price"], inplace=True)`
`print(df.head())`

```

symboling normalized-losses          make fuel-type aspiration \
0            3           122.0 alfa-romero      gas      std
1            3           122.0 alfa-romero      gas      std
2            1           122.0 alfa-romero      gas      std
3            2           164.0    audi      gas      std
4            2           164.0    audi      gas      std

  num-of-doors body-style drive-wheels engine-location wheelbase ... \
0        two   convertible         rwd       front     88.6 ...
1        two   convertible         rwd       front     88.6 ...
2        two   hatchback          rwd       front     94.5 ...
3       four    sedan            fwd       front     99.8 ...
4       four    sedan            4wd       front     99.4 ...

  engine-size fuel-system bore stroke compression-ratio horsepower \
0        130      mpfi  3.47   2.68          9.0     111.0
1        130      mpfi  3.47   2.68          9.0     111.0
2        152      mpfi  2.68   3.47          9.0     154.0
3        109      mpfi  3.19   3.40         10.0     102.0
4        136      mpfi  3.19   3.40          8.0     115.0

  peak-rpm city-mpg highway-mpg   price
0  5000.0       21        27 13495.0
1  5000.0       21        27 16500.0
2  5000.0       19        26 16500.0
3  5500.0       24        30 13950.0
4  5500.0       18        22 17450.0

```

[5 rows x 26 columns]

-> Replace missing values of "num-of-doors" with mode.

```
In [41]: df["num-of-doors"].fillna(df["num-of-doors"].mode()[0], inplace=True)
print(df.head())
```

```

symboling normalized-losses make fuel-type aspiration \
0            3      122.0 alfa-romero    gas     std
1            3      122.0 alfa-romero    gas     std
2            1      122.0 alfa-romero    gas     std
3            2      164.0      audi    gas     std
4            2      164.0      audi    gas     std

num-of-doors body-style drive-wheels engine-location wheelbase ... \
0           two   convertible        rwd      front    88.6 ...
1           two   convertible        rwd      front    88.6 ...
2           two   hatchback         rwd      front    94.5 ...
3          four    sedan           fwd      front    99.8 ...
4          four    sedan           4wd      front    99.4 ...

engine-size fuel-system bore stroke compression-ratio horsepower \
0          130      mpfi  3.47   2.68       9.0    111.0
1          130      mpfi  3.47   2.68       9.0    111.0
2          152      mpfi  2.68   3.47       9.0    154.0
3          109      mpfi  3.19   3.40      10.0    102.0
4          136      mpfi  3.19   3.40       8.0    115.0

peak-rpm city-mpg highway-mpg price
0    5000.0      21        27 13495.0
1    5000.0      21        27 16500.0
2    5000.0      19        26 16500.0
3    5500.0      24        30 13950.0
4    5500.0      18        22 17450.0

```

[5 rows x 26 columns]

C:\Users\rajva\AppData\Local\Temp\ipykernel_19968\2149171378.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["num-of-doors"].fillna(df["num-of-doors"].mode()[0], inplace=True)
```

-> Replace all other missing values with median.

```
In [39]: for col in df.columns:
    if df[col].isnull().any():
        df[col] = df[col].fillna(df[col].median())
print(df.head())
```

```

symboling    normalized-losses      make fuel-type aspiration \
0            3                  122.0 alfa-romero      gas      std
1            3                  122.0 alfa-romero      gas      std
2            1                  122.0 alfa-romero      gas      std
3            2                  164.0 audi           gas      std
4            2                  164.0 audi           gas      std

num-of-doors   body-style drive-wheels engine-location wheelbase ... \
0             two    convertible        rwd       front     88.6 ...
1             two    convertible        rwd       front     88.6 ...
2             two    hatchback         rwd       front     94.5 ...
3             four   sedan            fwd       front     99.8 ...
4             four   sedan            4wd       front     99.4 ...

engine-size   fuel-system bore stroke compression-ratio horsepower \
0            130      mpfi    3.47   2.68          9.0     111.0
1            130      mpfi    3.47   2.68          9.0     111.0
2            152      mpfi    2.68   3.47          9.0     154.0
3            109      mpfi    3.19   3.40         10.0     102.0
4            136      mpfi    3.19   3.40          8.0     115.0

peak-rpm city-mpg highway-mpg price
0      5000.0      21          27 13495.0
1      5000.0      21          27 16500.0
2      5000.0      19          26 16500.0
3      5500.0      24          30 13950.0
4      5500.0      18          22 17450.0

```

[5 rows x 26 columns]

4.Grouping, Sorting, and Aggregation

-> Group by "Fuel-type" and compute average price.

```
In [48]: df["price"] = pd.to_numeric(df["price"], errors='coerce')
avg_price_by_fuel = df.groupby("fuel-type")["price"].mean()
print("Average price by fuel-type:\n", avg_price_by_fuel)
```

```
Average price by fuel-type:
fuel-type
diesel    15838.15000
gas      12916.40884
Name: price, dtype: float64
```

-> In our dataset, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard We will need to apply data transformation to transform mpg into L/100km? The formula for unit conversion is L/100km = 235 / mpg

```
In [53]: df["city-L/100km"] = 235 / df["city-mpg"]
df["highway-L/100km"] = 235 / df["highway-mpg"]
print(df[["city-mpg", "city-L/100km", "highway-mpg", "highway-L/100km"]].head())
```

	city-mpg	city-L/100km	highway-mpg	highway-L/100km
0	21	11.190476	27	8.703704
1	21	11.190476	27	8.703704
2	19	12.368421	26	9.038462
3	24	9.791667	30	7.833333
4	18	13.055556	22	10.681818

-> Sort the DataFrame based on “price” in descending order.

```
In [57]: df_sorted = df.sort_values(by="price", ascending=False)
print(df_sorted[["make", "price"]].head())
```

	make	price
74	mercedes-benz	45400.0
16	bmw	41315.0
73	mercedes-benz	40960.0
128	porsche	37028.0
17	bmw	36880.0

```
In [ ]:
```

NumPy

1. NumPy Basic Ops

```
In [ ]: import numpy as np  
import pandas as pd
```

-> Create a 1D Array from 1 to 20.

```
In [186... array_1d = np.arange(1, 21)  
print("1D Array from 1 to 20:\n", array_1d)
```

1D Array from 1 to 20:
[1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20]

-> Create a 3×4 matrix of ones and reshape it to 4×3

```
In [189... mat_1 = np.ones((3,4))  
print("\n 3x4 Matrix of Ones: \n", mat_1)  
  
# Reshaping It to 4x3  
  
reshape = mat_1.reshape((4,3))  
print("\n 4x3 Reshaped Materix of Ones: \n", reshape)
```

3x4 Matrix of Ones:
[[1. 1. 1. 1.]
[1. 1. 1. 1.]
[1. 1. 1. 1.]]

4x3 Reshaped Materix of Ones:
[[1. 1. 1.]
[1. 1. 1.]
[1. 1. 1.]
[1. 1. 1.]]

-> Create a 5×5 identity matrix

```
In [194... i5 = np.eye(5)  
print("\n 5x5 Identity Matrix: \n", i5)
```

5x5 Identity Matrix:
[[1. 0. 0. 0. 0.]
[0. 1. 0. 0. 0.]
[0. 0. 1. 0. 0.]
[0. 0. 0. 1. 0.]
[0. 0. 0. 0. 1.]]

-> Generate 15 equally spaced numbers between 5 and 50.

In [199...]

```
e_nums = np.linspace(5,50,15)
print("\n 15 equally spaced nums between 5 and 15: \n", e_nums)
```

15 equally spaced nums between 5 and 15:
[5. 8.21428571 11.42857143 14.64285714 17.85714286 21.07142857
24.28571429 27.5 30.71428571 33.92857143 37.14285714 40.35714286
43.57142857 46.78571429 50.]

-> Generate a 4×4 matrix of random integers between 1 and 100.

In [201...]

```
r_nums = np.random.randint(1,101,size=(4,4))
print("\n 4x4 matrix of Random int between 1 & 100 \n",r_nums)
```

4x4 matrix of Random int between 1 & 100
[[48 17 75 76]
[47 94 95 48]
[36 68 89 44]
[93 58 61 56]]

2. NumPy Array Ops

-> Create a 3×3 matrix of random integers between 1 and 100

In [205...]

```
r_mat = np.random.randint(1,100,size=(3,3))
print("\n 3x3 Matrix of random nums \n",r_mat)
```

3x3 Matrix of random nums
[[34 34 99]
[14 14 29]
[37 85 1]]

-> Extract: First row, Second column, Center element.

In [208...]

```
fr = r_mat[0]
print("\n First Row: \n",fr)

sc = r_mat[:,1]
print("\n Second Column \n",sc)

se = r_mat[1,1]
print("\n Center Element \n",se)
```

First Row:

```
[34 34 99]
```

Second Column

```
[34 14 85]
```

Center Element

```
14
```

-> Replace all values greater than 50 in a matrix with 999.

```
In [211...]  
r_mat_c = r_mat.copy()  
r_mat_c[r_mat_c > 50] = 999  
print("\n Replaced Matrix \n",r_mat_c)
```

Replaced Matrix

```
[[ 34  34 999]  
[ 14  14  29]  
[ 37 999    1]]
```

-> Multiply a 1D array of size 5 (random integers between 1 to 10) by 10 using broadcasting.

```
In [214...]  
arr_1d = np.random.randint(1,11,size=5)  
multi_arr = arr_1d * 10  
print("\n Original 1D Array: \n",arr_1d)  
print("\n Multiplied Array: \n",multi_arr)
```

Original 1D Array:

```
[10  9  3  8  6]
```

Multiplied Array:

```
[100 90 30 80 60]
```

3. Math & Stats Ops

-> Create a 3×3 matrix of random integers between 1 and 100

```
In [225...]  
r_mat = np.random.randint(1,100,size=(3,3))  
print("\n 3x3 Matrix of random nums \n",r_mat)
```

3x3 Matrix of random nums

```
[[24  1 95]  
[79 17 61]  
[99 93 46]]
```

-> Compute sum, mean, median, std, var, min, and max of the above array.

```
In [230...]  
print("\nStatistics:")  
print("Sum:", np.sum(r_mat))  
print("Mean:", np.mean(r_mat))
```

```
print("Median:", np.median(r_mat))
print("Standard Deviation:", np.std(r_mat))
print("Variance:", np.var(r_mat))
print("Minimum:", np.min(r_mat))
print("Maximum:", np.max(r_mat))
```

Statistics:
Sum: 515
Mean: 57.22222222222222
Median: 61.0
Standard Deviation: 34.880041341010724
Variance: 1216.6172839506173
Minimum: 1
Maximum: 99

-> Normalize a 1D array of size 5 (random integers between 1 to 10) to scale values between 0 and 1.

```
In [234...]
arr_1d = np.random.randint(1,11,size=5)
print("\n 1D Array: \n",arr_1d)
min = np.min(arr_1d)
max = np.max(arr_1d)
norm_arr = ((arr_1d - min) / (max - min) )
print("\n Normalized Array: \n",norm_arr)
```

1D Array:
[8 10 5 1 4]

Normalized Array:
[0.77777778 1. 0.44444444 0. 0.33333333]

4. NumPy Matrix Ops & Algebra

-> Generate Following two NumPy matrix import numpy as np

A = ([4, 2], [1, 3])

B = ([2, 0], [1, 5])

```
In [240...]
A = np.array( [[4, 2], [1, 3]] )
B = np.array( [[2, 0], [1, 5]] )

print("\n A: \n",A)
print("\n B: \n",B)
```

```
A:  
[[4 2]  
[1 3]]
```

```
B:  
[[2 0]  
[1 5]]
```

-> Matrix Multiplication

```
In [252...]: mat_multi = np.matmul(A,B)  
print("\n Multiplication of A & B: \n",mat_multi)
```

```
Multiplication of A & B:  
[[10 10]  
[ 5 15]]
```

-> dot product of A and B.

```
In [254...]: dot = np.dot(A,B)  
print("\n dot product of A & B.: \n",dot)
```

```
dot product of A & B.:  
[[10 10]  
[ 5 15]]
```

-> Element wise addition/ subtraction/ multiplications/ division of A and B

```
In [250...]: add = A + B  
sub = A - B  
multi = A * B  
div = A / B  
  
print("\nElement-wise Addition:\n", add)  
print("Element-wise Subtraction:\n", sub)  
print("Element-wise Multiplication:\n", multi)  
print("Element-wise Division:\n", div)
```

```
Element-wise Addition:
```

```
[[6 2]  
[2 8]]
```

```
Element-wise Subtraction:
```

```
[[ 2  2]  
[ 0 -2]]
```

```
Element-wise Multiplication:
```

```
[[ 8  0]  
[ 1 15]]
```

```
Element-wise Division:
```

```
[[2. inf]  
[1. 0.6]]
```

```
C:\Users\rajva\AppData\Local\Temp\ipykernel_18728\3905969198.py:4: RuntimeWarning: divide by zero encountered in divide  
div = A / B
```

-> Transpose of A

```
In [277...]: trans_A = A.T  
print("\n Transpose of A: \n",trans_A)
```

```
Transpose of A:  
[[4 1]  
[2 3]]
```

-> Determinant of A

```
In [271...]: det_A = np.linalg.det(A)  
print("\n Determinant of A: \n",det_A)
```

```
Determinant of A:  
10.00000000000002
```

-> Inverse of A

```
In [269...]: if det_A != 0:  
    inv_A = np.linalg.inv(A)  
    print("\nInverse of A:\n", inv_A)  
else:  
    print("\nMatrix A cannot be inverted.")
```

```
Inverse of A:  
[[ 0.3 -0.2]  
[-0.1  0.4]]
```

-> Eigenvalues and Eigenvectors of A

```
In [264...]: eigenvalues= np.linalg.eig(A)  
eigenvectors = np.linalg.eig(A)  
print("\nEigenvalues of A:\n", eigenvalues)  
print("\nEigenvectors of A:\n", eigenvectors)
```

```
Eigenvalues of A:  
EigResult(eigenvalues=array([5., 2.]), eigenvectors=array([[ 0.89442719, -0.70710678],  
[ 0.4472136 ,  0.70710678]]))
```

```
Eigenvectors of A:  
EigResult(eigenvalues=array([5., 2.]), eigenvectors=array([[ 0.89442719, -0.70710678],  
[ 0.4472136 ,  0.70710678]]))
```

-> Solve System of Equations:

$$2x + y = 8$$

$$3x + 4y = 18$$

```
In [261...]: coeff = np.array([[2,1],[3,4]])
cons = np.array([8,18])

sol = np.linalg.solve(coeff, cons)
print("\n Solution of the Equations x & y : ",sol)

Solution of the Equations x & y : [2.8 2.4]
```

Pandas Exercise

1. Series & DataFrame Basics

-> Given the following list of marks:

[78, 85, 92, 70, 66]

Create a Pandas Series and assign the following student names as indices:

['Amit', 'Bhavna', 'Chetan', 'Divya', 'Esha']

Display the Series

```
In [ ]: marks = [78, 85, 92, 70, 66]
names = ['Amit', 'Bhavna', 'Chetan', 'Divya', 'Esha']

marks_series = pd.Series(data=marks, index=names)
print("\nPandas Series : \n",marks_series)
```

-> Create a Pandas DataFrame and display:

The full DataFrame

The column names

The shape of the DataFrame

```
In [285...]: # Create DataFrame
data = {
    'Name': ['Amit', 'Bhavna', 'Chetan', 'Divya', 'Esha'],
    'Gender': ['Male', 'Female', 'Male', 'Female', 'Female'],
    'Math': [78, 85, 92, 70, 66],
    'Science': [88, 79, 95, 72, 60]
}

df = pd.DataFrame(data)

print("\nFull DataFrame:\n")
print(df)
```

```

# Display column names
print("\nColumn Names:\n")
print(df.columns.tolist())

# Display shape (rows, columns)
print("\nShape of the DataFrame:\n")
print(df.shape)

```

Full DataFrame:

	Name	Gender	Math	Science
0	Amit	Male	78	88
1	Bhavna	Female	85	79
2	Chetan	Male	92	95
3	Divya	Female	70	72
4	Esha	Female	66	60

Column Names:

```
[ 'Name', 'Gender', 'Math', 'Science' ]
```

Shape of the DataFrame:

```
(5, 4)
```

2.Data Exploration

-> Load the dataset from this url
["https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data"](https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data)

In [7]:

```

import pandas as pd
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data"
df = pd.read_csv(url, header=None)
print(df.head())

```

0	1	2	3	4	5	6	7	8	9	...	\		
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...		
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...		
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...		
3	2	164		audi	gas	std	four		sedan	fwd	front	99.8	...
4	2	164		audi	gas	std	four		sedan	4wd	front	99.4	...
	16	17	18	19	20	21	22	23	24	25			
0	130	mpfi	3.47	2.68	9.0	111	5000	21	27	13495			
1	130	mpfi	3.47	2.68	9.0	111	5000	21	27	16500			
2	152	mpfi	2.68	3.47	9.0	154	5000	19	26	16500			
3	109	mpfi	3.19	3.40	10.0	102	5500	24	30	13950			
4	136	mpfi	3.19	3.40	8.0	115	5500	18	22	17450			

[5 rows x 26 columns]

-> Assign names of columns:

```
["symboling","normalized-losses","make","fuel-type","aspiration", "num-ofdoors","body-style", "drive-wheels","engine-ocation","wheelbase","length","width","height","curb-weight","engine-type","num-of-cylinders","engine-size","fuel-system","bore","stroke","compression-ratio","horsepower","peak-rpm","city-mpg","highway-mpg","price"]
```

```
In [296... df.columns = ["symboling","normalized-losses","make","fuel-type","aspiration", "num-of-doors","body-style", "drive-wheels","engine-location","wheel-base","length","width","height","curb-weight","engine-type","num-of-cylinders","engine-size","fuel-system","bore","stroke","compression-ratio","horsepower","peak-rpm","city-mpg","highway-mpg","price"]
print(df.head())
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height	curb-weight	engine-type	num-of-cylinders	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg	price
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...				mpfi	3.47	2.68	9.0	111	5000	21					
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...				mpfi	3.47	2.68	9.0	111	5000	21					
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...				mpfi	2.68	3.47	9.0	154	5000	19					
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...				mpfi	3.19	3.40	10.0	102	5500	24					
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...				mpfi	3.19	3.40	8.0	115	5500	18					

[5 rows x 26 columns]

-> Display .shape, .columns, .info(), and .describe()

```
In [300... print(df.shape)
print(df.columns)
print(df.info())
print(df.describe())
```

```
(205, 26)
Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
       'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
       'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
       'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
       'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
       'highway-mpg', 'price'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   symboling         205 non-null    int64  
 1   normalized-losses 205 non-null    object  
 2   make              205 non-null    object  
 3   fuel-type         205 non-null    object  
 4   aspiration        205 non-null    object  
 5   num-of-doors      205 non-null    object  
 6   body-style        205 non-null    object  
 7   drive-wheels      205 non-null    object  
 8   engine-location    205 non-null    object  
 9   wheel-base        205 non-null    float64 
 10  length             205 non-null    float64 
 11  width              205 non-null    float64 
 12  height             205 non-null    float64 
 13  curb-weight        205 non-null    int64  
 14  engine-type        205 non-null    object  
 15  num-of-cylinders   205 non-null    object  
 16  engine-size        205 non-null    int64  
 17  fuel-system        205 non-null    object  
 18  bore               205 non-null    object  
 19  stroke             205 non-null    object  
 20  compression-ratio   205 non-null    float64 
 21  horsepower          205 non-null    object  
 22  peak-rpm            205 non-null    object  
 23  city-mpg            205 non-null    int64  
 24  highway-mpg          205 non-null    int64  
 25  price              205 non-null    object  
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB
None
      symboling  wheel-base    length     width     height \
count  205.000000  205.000000  205.000000  205.000000  205.000000
mean    0.834146   98.756585  174.049268  65.907805   53.724878
std     1.245307   6.021776  12.337289  2.145204   2.443522
min    -2.000000   86.600000  141.100000  60.300000  47.800000
25%    0.000000   94.500000  166.300000  64.100000  52.000000
50%    1.000000   97.000000  173.200000  65.500000  54.100000
75%    2.000000  102.400000  183.100000  66.900000  55.500000
max    3.000000  120.900000  208.100000  72.300000  59.800000

      curb-weight  engine-size  compression-ratio  city-mpg  highway-mpg
count  205.000000  205.000000      205.000000  205.000000  205.000000
mean   2555.565854  126.907317      10.142537  25.219512  30.751220
std    520.680204   41.642693      3.972040   6.542142  6.886443
```

```
min    1488.000000    61.000000    7.000000    13.000000    16.000000
25%   2145.000000    97.000000    8.600000    19.000000    25.000000
50%   2414.000000   120.000000    9.000000    24.000000    30.000000
75%   2935.000000   141.000000    9.400000    30.000000    34.000000
max    4066.000000   326.000000   23.000000    49.000000    54.000000
```

-> Display only "width", "height", "curb-weight", "engine-type" columns.

```
In [303...]: print(df[["width", "height", "curb-weight", "engine-type"]])
```

```
   width  height  curb-weight engine-type
0     64.1    48.8       2548      dohc
1     64.1    48.8       2548      dohc
2     65.5    52.4       2823      ohcv
3     66.2    54.3       2337      ohc
4     66.4    54.3       2824      ohc
..     ...
200    68.9    55.5       2952      ohc
201    68.8    55.5       3049      ohc
202    68.9    55.5       3012      ohcv
203    68.9    55.5       3217      ohc
204    68.9    55.5       3062      ohc
```

[205 rows x 4 columns]

-> Display car details which have num-of-doors = four

```
In [307...]: print(df[df["num-of-doors"] == "four"])
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	
6	1	158	audi	gas	std	four	
7	1	?	audi	gas	std	four	
8	1	158	audi	gas	turbo	four	
..	
200	-1	95	volvo	gas	std	four	
201	-1	95	volvo	gas	turbo	four	
202	-1	95	volvo	gas	std	four	
203	-1	95	volvo	diesel	turbo	four	
204	-1	95	volvo	gas	turbo	four	
	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	
6	sedan	fwd	front	105.8	...	136	
7	wagon	fwd	front	105.8	...	136	
8	sedan	fwd	front	105.8	...	131	
..	
200	sedan	rwd	front	109.1	...	141	
201	sedan	rwd	front	109.1	...	141	
202	sedan	rwd	front	109.1	...	173	
203	sedan	rwd	front	109.1	...	145	
204	sedan	rwd	front	109.1	...	141	
	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	\
3	mpfi	3.19	3.40	10.0	102	5500	
4	mpfi	3.19	3.40	8.0	115	5500	
6	mpfi	3.19	3.40	8.5	110	5500	
7	mpfi	3.19	3.40	8.5	110	5500	
8	mpfi	3.13	3.40	8.3	140	5500	
..	
200	mpfi	3.78	3.15	9.5	114	5400	
201	mpfi	3.78	3.15	8.7	160	5300	
202	mpfi	3.58	2.87	8.8	134	5500	
203	idi	3.01	3.40	23.0	106	4800	
204	mpfi	3.78	3.15	9.5	114	5400	
	city-mpg	highway-mpg	price				
3	24	30	13950				
4	18	22	17450				
6	19	25	17710				
7	19	25	18920				
8	17	20	23875				
..				
200	23	28	16845				
201	19	25	19045				
202	18	23	21485				
203	26	27	22470				
204	19	25	22625				

[114 rows x 26 columns]

3. Missing Value Handlling

-> Missing value is represented by '?' in this dataset. Replace it with NULL.

```
In [311... df.replace('?', None , inplace=True)
print(df.head())
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	3	None	alfa-romero	gas	std	two	
1	3	None	alfa-romero	gas	std	two	
2	1	None	alfa-romero	gas	std	two	
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

-> Check how many missing values are there in each attribute.

```
In [316... print(df.isnull().sum())
```

```
symboling          0
normalized-losses 41
make              0
fuel-type         0
aspiration        0
num-of-doors      2
body-style        0
drive-wheels      0
engine-location   0
wheel-base        0
length            0
width             0
height            0
curb-weight       0
engine-type       0
num-of-cylinders 0
engine-size       0
fuel-system       0
bore               4
stroke             4
compression-ratio 0
horsepower         2
peak-rpm           2
city-mpg           0
highway-mpg        0
price              4
dtype: int64
```

-> Replace missing values of "normalized-losses", "stroke", "bore", "horsepower" with mean.

```
In [26]: for col in ["normalized-losses", "stroke", "bore", "horsepower"]:
    df[col].fillna(df[col].mean(), inplace=True)
print(df.head())
```

```

symboling    normalized-losses      make fuel-type aspiration \
0            3                  122.0 alfa-romero      gas      std
1            3                  122.0 alfa-romero      gas      std
2            1                  122.0 alfa-romero      gas      std
3            2                  164.0 audi           gas      std
4            2                  164.0 audi           gas      std

num-of-doors   body-style drive-wheels engine-location wheelbase ... \
0             two    convertible        rwd       front     88.6 ...
1             two    convertible        rwd       front     88.6 ...
2             two    hatchback         rwd       front     94.5 ...
3             four   sedan            fwd       front     99.8 ...
4             four   sedan            4wd       front     99.4 ...

engine-size   fuel-system bore stroke compression-ratio horsepower \
0            130      mpfi    3.47    2.68          9.0      111.0
1            130      mpfi    3.47    2.68          9.0      111.0
2            152      mpfi    2.68    3.47          9.0      154.0
3            109      mpfi    3.19    3.40         10.0      102.0
4            136      mpfi    3.19    3.40          8.0      115.0

peak-rpm city-mpg highway-mpg price
0      5000.0      21          27 13495.0
1      5000.0      21          27 16500.0
2      5000.0      19          26 16500.0
3      5500.0      24          30 13950.0
4      5500.0      18          22 17450.0

```

[5 rows x 26 columns]

C:\Users\rajva\AppData\Local\Temp\ipykernel_19968\497043668.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an `inplace` method.

The behavior will change in pandas 3.0. This `inplace` method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation `inplace` on the original object.

```
df[col].fillna(df[col].mean(), inplace=True)
```

-> Drop all the rows which has missing value in attribute "price".

```
In [31]: df.dropna(subset=["price"], inplace=True)
print(df.head())
```

```

symboling normalized-losses          make fuel-type aspiration \
0            3           122.0 alfa-romero      gas      std
1            3           122.0 alfa-romero      gas      std
2            1           122.0 alfa-romero      gas      std
3            2           164.0     audi      gas      std
4            2           164.0     audi      gas      std

num-of-doors body-style drive-wheels engine-location wheelbase ... \
0            two   convertible        rwd       front    88.6 ...
1            two   convertible        rwd       front    88.6 ...
2            two   hatchback         rwd       front    94.5 ...
3            four    sedan           fwd       front    99.8 ...
4            four    sedan           4wd       front    99.4 ...

engine-size fuel-system bore stroke compression-ratio horsepower \
0           130      mpfi   3.47   2.68          9.0    111.0
1           130      mpfi   3.47   2.68          9.0    111.0
2           152      mpfi   2.68   3.47          9.0    154.0
3           109      mpfi   3.19   3.40         10.0    102.0
4           136      mpfi   3.19   3.40          8.0    115.0

peak-rpm city-mpg highway-mpg   price
0      5000.0      21        27 13495.0
1      5000.0      21        27 16500.0
2      5000.0      19        26 16500.0
3      5500.0      24        30 13950.0
4      5500.0      18        22 17450.0

```

[5 rows x 26 columns]

-> Replace missing values of “num-of-doors” with mode.

```
In [41]: df["num-of-doors"].fillna(df["num-of-doors"].mode()[0], inplace=True)
print(df.head())
```

```

symboling normalized-losses          make fuel-type aspiration \
0            3           122.0 alfa-romero      gas      std
1            3           122.0 alfa-romero      gas      std
2            1           122.0 alfa-romero      gas      std
3            2           164.0     audi      gas      std
4            2           164.0     audi      gas      std

num-of-doors body-style drive-wheels engine-location wheelbase ... \
0            two   convertible        rwd       front    88.6 ...
1            two   convertible        rwd       front    88.6 ...
2            two   hatchback         rwd       front   94.5 ...
3            four    sedan           fwd       front   99.8 ...
4            four    sedan           4wd       front   99.4 ...

engine-size fuel-system bore stroke compression-ratio horsepower \
0           130      mpfi   3.47    2.68          9.0    111.0
1           130      mpfi   3.47    2.68          9.0    111.0
2           152      mpfi   2.68    3.47          9.0    154.0
3           109      mpfi   3.19    3.40         10.0    102.0
4           136      mpfi   3.19    3.40          8.0    115.0

peak-rpm city-mpg highway-mpg price
0      5000.0      21        27 13495.0
1      5000.0      21        27 16500.0
2      5000.0      19        26 16500.0
3      5500.0      24        30 13950.0
4      5500.0      18        22 17450.0

```

[5 rows x 26 columns]

C:\Users\rajva\AppData\Local\Temp\ipykernel_19968\2149171378.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an `inplace` method.

The behavior will change in pandas 3.0. This `inplace` method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation `inplace` on the original object.

```
df["num-of-doors"].fillna(df["num-of-doors"].mode()[0], inplace=True)
```

-> Replace all other missing values with median.

```
In [39]: for col in df.columns:
    if df[col].isnull().any():
        df[col] = df[col].fillna(df[col].median())
print(df.head())
```

```

symboling    normalized-losses      make fuel-type aspiration \
0            3                  122.0 alfa-romero      gas      std
1            3                  122.0 alfa-romero      gas      std
2            1                  122.0 alfa-romero      gas      std
3            2                  164.0 audi           gas      std
4            2                  164.0 audi           gas      std

num-of-doors   body-style drive-wheels engine-location wheelbase ... \
0             two    convertible        rwd       front     88.6 ...
1             two    convertible        rwd       front     88.6 ...
2             two    hatchback         rwd       front     94.5 ...
3             four   sedan            fwd       front     99.8 ...
4             four   sedan            4wd       front     99.4 ...

engine-size   fuel-system bore stroke compression-ratio horsepower \
0            130      mpfi    3.47   2.68          9.0      111.0
1            130      mpfi    3.47   2.68          9.0      111.0
2            152      mpfi    2.68   3.47          9.0      154.0
3            109      mpfi    3.19   3.40         10.0      102.0
4            136      mpfi    3.19   3.40          8.0      115.0

peak-rpm city-mpg highway-mpg price
0      5000.0      21          27 13495.0
1      5000.0      21          27 16500.0
2      5000.0      19          26 16500.0
3      5500.0      24          30 13950.0
4      5500.0      18          22 17450.0

```

[5 rows x 26 columns]

4.Grouping, Sorting, and Aggregation

-> Group by “Fuel-type” and compute average price.

```
In [48]: df["price"] = pd.to_numeric(df["price"], errors='coerce')
avg_price_by_fuel = df.groupby("fuel-type")["price"].mean()
print("Average price by fuel-type:\n", avg_price_by_fuel)
```

```
Average price by fuel-type:
fuel-type
diesel    15838.15000
gas      12916.40884
Name: price, dtype: float64
```

-> In our dataset, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard We will need to apply data transformation to transform mpg into L/100km? The formula for unit conversion is L/100km = 235 / mpg

```
In [53]:
```

```
df["city-L/100km"] = 235 / df["city-mpg"]
df["highway-L/100km"] = 235 / df["highway-mpg"]
print(df[["city-mpg", "city-L/100km", "highway-mpg", "highway-L/100km"]].head())
```

	city-mpg	city-L/100km	highway-mpg	highway-L/100km
0	21	11.190476	27	8.703704
1	21	11.190476	27	8.703704
2	19	12.368421	26	9.038462
3	24	9.791667	30	7.833333
4	18	13.055556	22	10.681818

-> Sort the DataFrame based on “price” in descending order.

```
In [57]:
```

```
df_sorted = df.sort_values(by="price", ascending=False)
print(df_sorted[["make", "price"]].head())
```

	make	price
74	mercedes-benz	45400.0
16	bmw	41315.0
73	mercedes-benz	40960.0
128	porsche	37028.0
17	bmw	36880.0

```
In [ ]:
```